

# Introduction to JPA and Hibernate

Held as part of the lecture series on  
Web Engineering at Vienna University of Technology

May 2014



Business Informatics Group

Vienna University of Technology

# Web Engineering

## Introduction to JPA and Hibernate



**Philipp Liegl**

**Business Informatics Group**

Institute of Software Technology and Interactive Systems  
Vienna University of Technology

Favoritenstraße 9-11/188-3, 1040 Vienna, Austria  
phone: +43 (1) 58801-18804 (secretary), fax: +43 (1) 58801-18896  
office@big.tuwien.ac.at, www.big.tuwien.ac.at

# Outline of today's talk

---

- JDBC
- JPA/Hibernate
  - Relationships
  - Persistence Context/Persistence Unit
  - Entity Manager
  - JPQL
  - Hibernate Criteria API

Accompanying examples


<https://github.com/pliegl/we2014/tree/master/jpa-sample>

# Motivation

## N-Tier Architectures

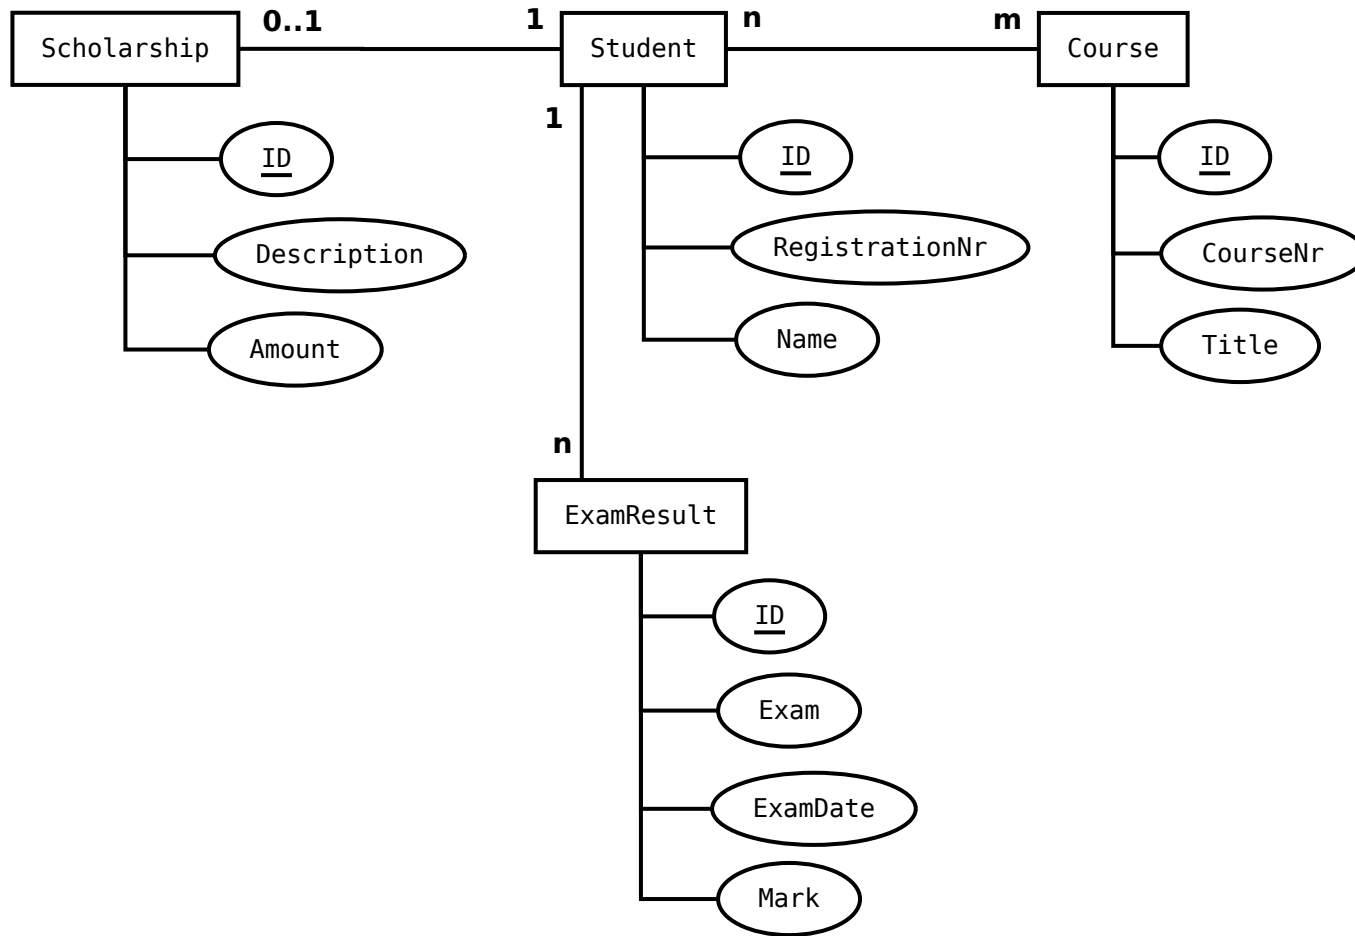
### ■ Layers of an information system

- Presentation layer
  - Communication interface to external entities
  - “View” in the model-view-controller
- Application logic layer (service layer)
  - Implements operations requested by clients through the presentation layer
  - Represents the “business logic”
- Resource management layer (persistence layer)
  - Deals with different data sources of an information system
  - Responsible for storing and retrieving data



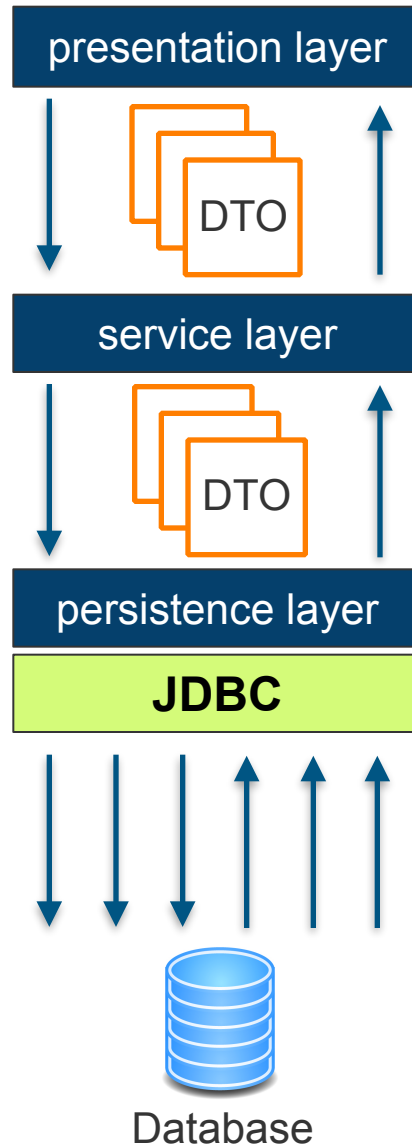
# Motivation

## Accompanying model



# Motivation

## Traditional persistence with JDBC



# Motivation

## JDBC - Java Database Connectivity

---

- Used to access relational databases from Java programs
- First version released 1996
- Ability to
  - Establish a connection to a database
  - Execute an SQL statement and return results
  - Create parameterized queries
  - Manage database transactions
- Basic Steps
  - Load driver or obtain an already defined data source
  - Establish connection using a JDBC URL
  - Create an **SQL statement** and execute SQL statement
  - If present, process results present in **result sets**
  - Close database resources
  - Commit or rollback transaction, if necessary



# JDBC example

Insert an entry

```
Connection conn = null;
PreparedStatement stmt = null;

try {
    conn = connection();
    stmt = conn.prepareStatement( "INSERT INTO student VALUES(?, ?, ?)" );
    stmt.setInt( 1, student.getId() );
    stmt.setString( 2, student.getMatrNr() );
    stmt.setString( 3, student.getName() );
    stmt.executeUpdate();
    stmt.close();
} catch (Exception e) {
    e.printStackTrace();
} finally {
    if (stmt != null) {
        stmt.close();
    }
    if (conn != null) {
        conn.close();
    }
}
```



# JDBC example

Retrieve an entry

```
Connection conn = null;
PreparedStatement stmt = null;
ResultSet rs = null;
try {
    conn = connection();
    stmt = conn.prepareStatement( "SELECT id, matnr, name FROM student
                                  WHERE id=?" );

    stmt.setInt( 1, id );
    rs = stmt.executeQuery();
    rs.next();

    Student student = new Student();
    student.setId( rs.getInt( 1 ) );
    student.setMatrNr( rs.getString( 2 ) );
    student.setName( rs.getString( 3 ) );

    rs.close();
    stmt.close();
    return student;
} catch (Exception e) {
    e.printStackTrace();
} finally {
    ...
}
```

# JDBC

## Drawbacks

---

- Verbose JDBC boilerplate code for the various CRUD actions
- Manual mapping of JDBC result sets to the respective Java POJOs
  - Imagine 40 different database tables with 20 attributes each...
- Manual synchronization of Java code in case of database schema changes (e.g., a new field is added to a database table)
  - Manual adaptation of the entire related JDBC Java code necessary

# Object Relational Mapping

## Reasons for using ORM

---

- In an application we want to focus on **business concepts**, not on the relational database structure
- Abstract from the “by-hand” communication with the DB (e.g., via JDBC)
- Allow for an automatic synchronization between Java Objects and the underlying database
- Portability
  - ORM should be mostly DB independent (with the exception of some types of features, such as identifier generation)
  - Query abstractions using e.g. JPQL or HQL - the vendor specific SQL is auto-generated
- Performance
  - Object and query caching is automatically done by the ORM

# Java Persistence API (JPA)

## Introduction

---

- Specification for the management of persistence and object/relational mapping with Java
  - Persistence: Data objects shall outlive the JVM app
- Objective: provide an object/relational mapping facility for Java developers using a Java domain model and a relational database
  - Map Java POJOs to relational databases (which are one type of persistence)
- Standardized under the Java Community Process Program with contributions from Hibernate, TopLink, JDO, and the EJB community
- Hibernate: Full JPA implementation with additional “native” features, e.g.,
  - HQL (Hibernate Query Language) - similar to JPQL, but with some extensions
  - Criteria API
  - Used version in this course: Hibernate 4.2.12.Final (supports JPA 2.0)



# Caveats

## ORM and JPA

---

- With JPA/Hibernate lots of “magic” is done under the hood, e.g., SQL-DDL is automatically generated
- Know the database basics first (e.g., from a data engineering course), in order to fully understand what JPA is doing under the hood
- After annotating the classes and running the application check the resulting SQL-DDL (e.g., using the database explorer in IntelliJ or Eclipse)
- When executing SQL-Queries using JPA/Hibernate use the “show SQL queries” feature during development, in order to see what kind of queries are actually executed
  - **Set** `<property name="hibernate.show_sql" value="true" />` in `persistence.xml`

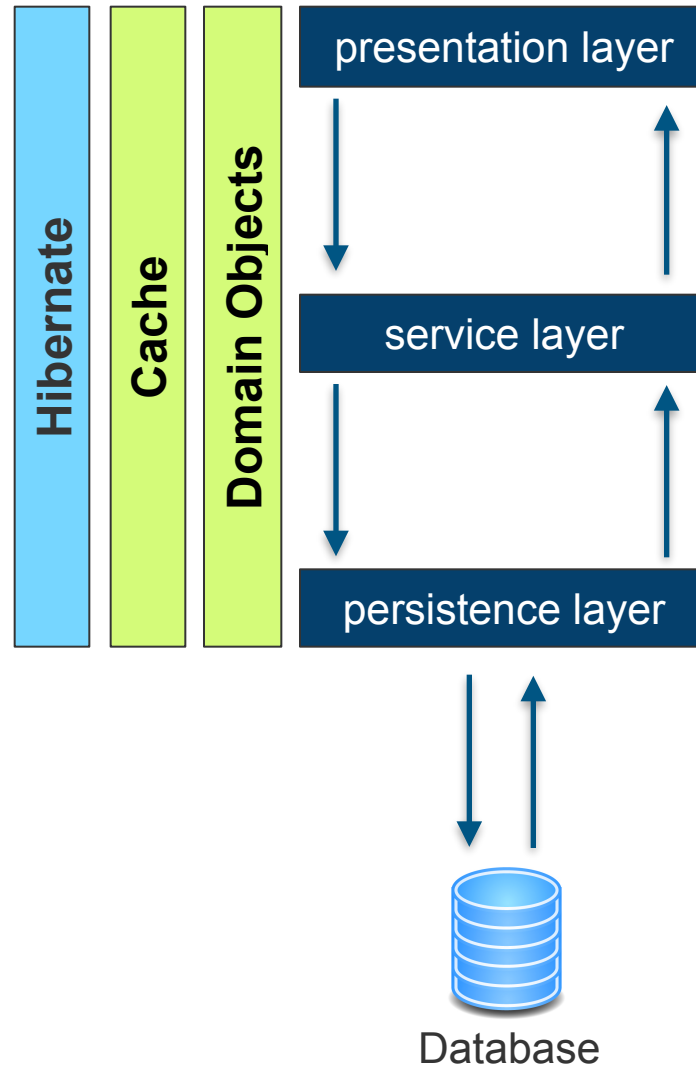
# Persistent Entities

## Basics

---

- Are POJOs (Plain Old Java Objects)
- Lightweight persistent domain object
- Typically represent a table in a relational database
- Each entity instance corresponds to one row in that table
- Have a persistent identity
- May have both, persistent and transient (non-persistent) state
  - Simple types (primitive data types, wrappers, enums)
  - Composite types (e.g., Address)
  - Non-persistent state (using identifier `transient` or `@Transient` annotation)

# Persistence with Hibernate



# Simple Mapping

Enhance Java domain classes with JPA annotations

```
@Entity
public class ExamResult {

    @Id
    private Long id;

    @Column(name = "pruefungDatum")
    @Temporal(TemporalType.DATE)
    private Date examDate;


    private int mark;

    @Transient
    private String examLocation;

    ...
}
```



ExamResult		
id	pruefungDatum	mark



## Important annotations

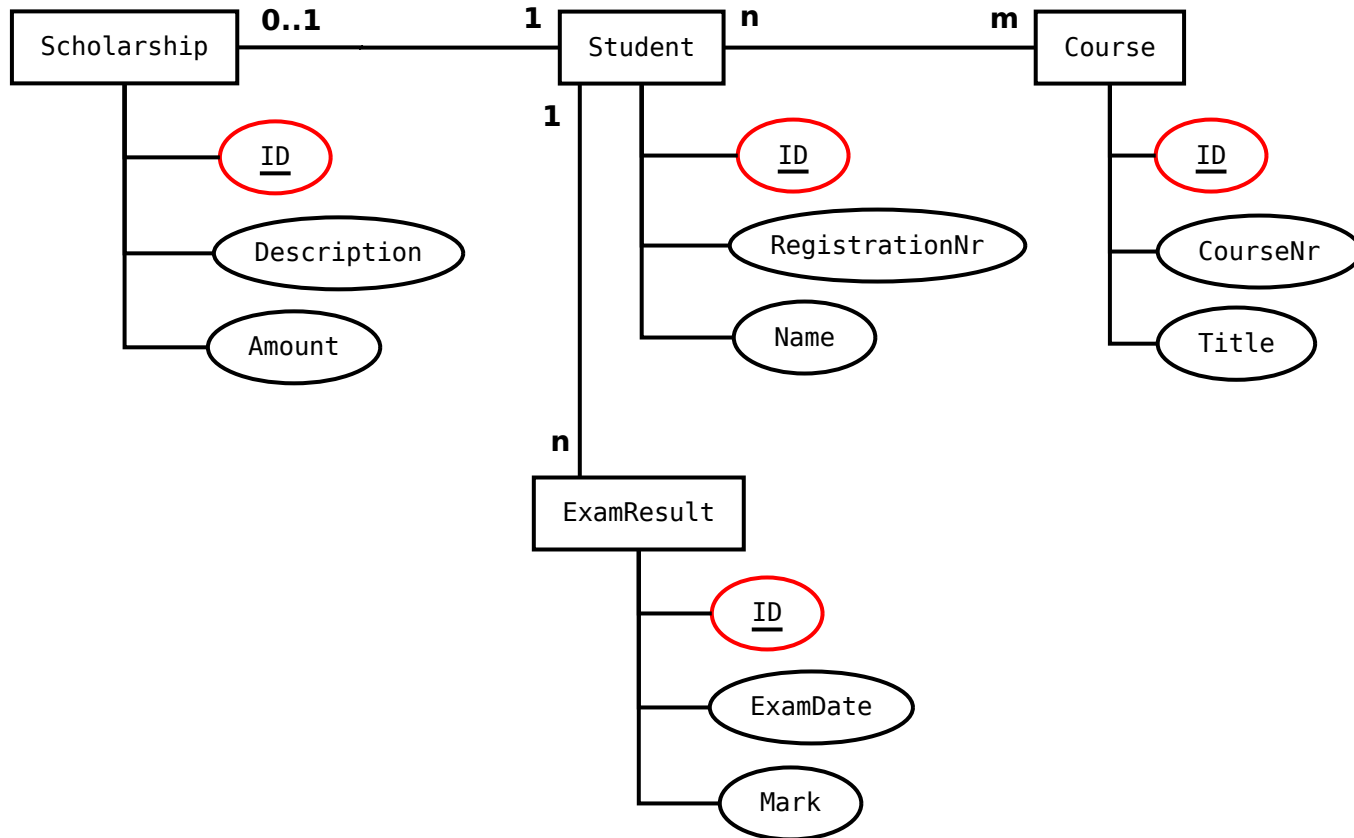
<b>@Entity</b>	Specifies that the class is an entity
<b>@Id</b>	Specifies the primary key of an entity
<b>@Temporal</b>	Must be specified for fields of type java.util.Date and java.util.Calendar
<b>@TemporalType</b>	Type used to indicate a specific mapping of java.util.Date or java.util.Calendar. Allowed values: <ul style="list-style-type: none"><li>- DATE</li><li>- TIME</li><li>- TIMESTAMP</li></ul>
<b>@Transient</b>	Specifies that the field is not persistent





# Simple Mapping

## Inheritance



# Simple Mapping

## Inheritance

```
@MappedSuperclass
public class BaseEntity {

    @Id
    @GeneratedValue(
        strategy =
        GenerationType.AUTO)
    protected Long id;

    public Long getId() {
        return id;
    }
}
```

### Important annotations

<b>@MappedSuperclass</b>	Designates a class whose mapping information is applied to the entities that inherit from it. A mapped superclass has no separate table defined for it.
<b>@GeneratedValue</b>	The <code>GeneratedValue</code> annotation may be applied to a primary key property or field of an entity or mapped superclass in conjunction with the <code>Id</code> annotation.

```
@Entity
public class ExamResult extends BaseEntity {

    @Column(name = "pruefungsDatum")
    @Temporal(TemporalType.DATE)
    private Date examDate;

    private int mark;

    @Transient
    private String examLocation;

    //Getter and setters omitted

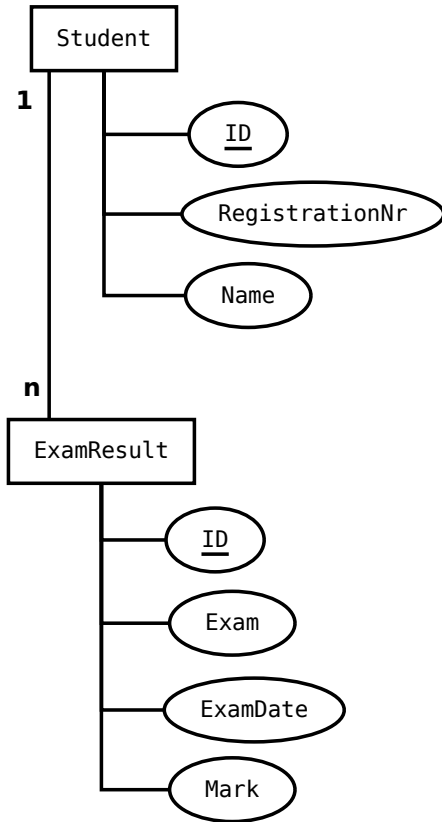
}
```



ExamResult		
id	pruefungsDatum	mark



# Object-oriented vs. SQL



OO: **Student owns the ExamResults**  
**Usually:** no ExamResult without a student

```
public class Student extends BaseEntity {

    private String registrationNumber;
    private String name;
    private List<ExamResult> examResults;
    ...
}
```

Does not exist in the DB, but is simulated using an SQL query. JPA takes care of that.

```
public class ExamResult extends BaseEntity {

    private Date examDate;
    private String exam;
    private int mark;
    private Student student;
    ...
}
```

SQL:

- ExamResult contains a foreign key to the Student it belongs to
- The ExamResult owns (contains) the connection
- This is opposite to the OO perspective

# Entity relationships

---

**One-to-one, one-to-many, many-to-many, many-to-one** relationships among entities

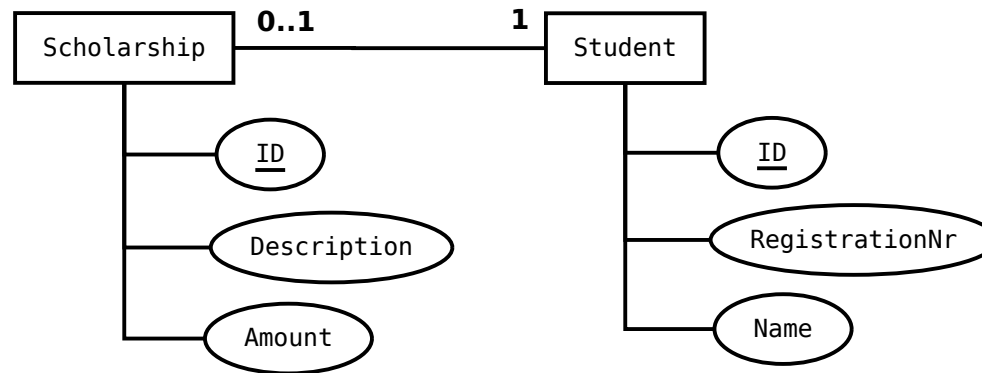
- bi-directional or uni-directional
- Support for different Collection types, e.g., List, Set, Map, etc.

Need to specify the **owning** side in relationships

- Owning side table has the foreign key
- OneToOne relationship - the side where the foreign key is specified
- OneToMany, ManyToOne - the “many” side

# Relationship mapping

Example using a unidirectional mapping




## Four different options:

1. Using an embedded table, where Scholarship is the embedded table. (see example)
2. Scholarship and Student are separate tables. The primary key of Scholarship has a foreign key constraint on the primary key of the “owning” Student (using `@PrimaryKeyJoinColumn` annotation)
3. Scholarship and Student are separate tables. Student holds a foreign key which references the primary key of Scholarship. The foreign key has a unique constraint.
4. Scholarship and Student are separate tables. Scholarship holds a foreign key which references the primary key of Student. The foreign key has a unique constraint. (see example)

# Relationship mapping

## Unidirectional OneToOne using an embedded table



### Important annotations

#### @Embedded

Defines a persistent field or property of an entity whose value is an instance of an embeddable class. The embeddable class must be annotated as `Embeddable`.

#### @Embeddable

Defines a class whose instances are stored as an intrinsic part of an owning entity and share the identity of the entity. Each of the persistent properties or fields of the embedded object is mapped to the database table for the entity.

#### @Entity

```
public class EmbeddedStudent extends BaseEntity {

    @Column(name = "matrikelNummer", unique = true)
    private String registrationNumber;

    private String name;

    @Embedded
    private EmbeddedScholarship scholarship;

    @Transient
    private DateTime loginTime;

    ...
}
```

#### @Embeddable

```
public class EmbeddedScholarship {

    private String description;

    private Integer amount;

}
```

Might be an issue with legacy databases, where the DB schema already exists and must not be altered.

# Relationship mapping

## Unidirectional OneToOne using an embedded table - resulting SQL DDL

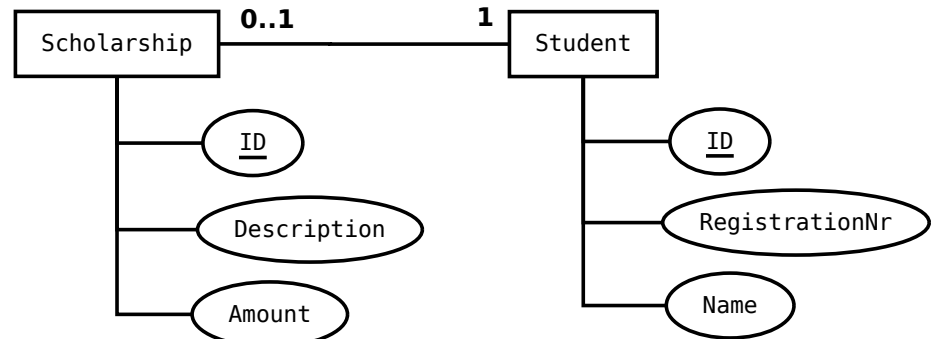
### @Entity

```
public class EmbeddedStudent extends BaseEntity {  
  
    @Column(name = "matrikelNummer", unique = true)  
    private String registrationNumber;  
  
    private String name;  
  
    @Embedded  
    private EmbeddedScholarship scholarship;  
  
    @Transient  
    private DateTime loginTime;  
  
    ...  
}
```

### @Embeddable


```
public class EmbeddedScholarship {  
  
    private String description;  
  
    private Integer amount;  
  
}
```

```
CREATE TABLE EMBEDDEDSTUDENT  
(  
    ID BIGINT PRIMARY KEY NOT NULL,  
    NAME VARCHAR(255),  
    MATRIKELNUMMER VARCHAR(255),  
    AMOUNT INTEGER,  
    DESCRIPTION VARCHAR(255)  
);  
CREATE UNIQUE INDEX anIndexName  
ON EMBEDDEDSTUDENT ( MATRIKELNUMMER );
```



# Relationship mapping

## Bidirectional OneToOne using foreign key



The “non-owning” side

### @Entity

```
public class Student extends BaseEntity {
```

```
    @Column(name = "matrikelNummer",
            unique = true)
    private String registrationNumber;
```

```
    private String name;
```

```
    @OneToOne(fetch = FetchType.LAZY,
            cascade = CascadeType.ALL,
            mappedBy="grantedTo")
    private Scholarship scholarship;
```

### @Transient

```
    private DateTime loginTime;
```

```
...
}
```

### Important annotations

#### @OneToOne

Defines a single-valued association to another entity that has one-to-one multiplicity.

#### @FetchType

LAZY = do not load referenced entity, until it is accessed for the first time  
EAGER = load referenced entity immediately

#### @CascadeType

Defines the set of cascadable operations that are propagated to the associated entity. ALL is equivalent to cascade={PERSIST, MERGE, REMOVE, REFRESH, DETACH}

#### mappedBy

References the field that “owns” the relationship in the referenced entity. Required unless the relationship is unidirectional.



# Relationship mapping

## Bidirectional OneToOne using foreign key cont'd

The “owning” side

```
@Entity
public class Scholarship extends BaseEntity {

    private String description;

    private Integer amount;

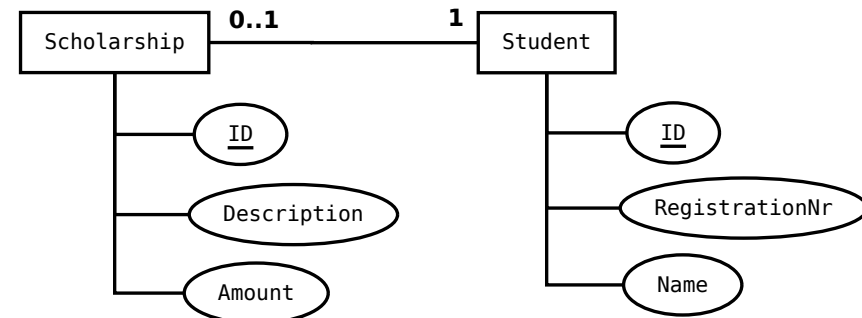
    @JoinColumn(name="student_id", unique=true)
    @OneToOne
    private Student grantedTo;

    ...
}
```

### Important annotations

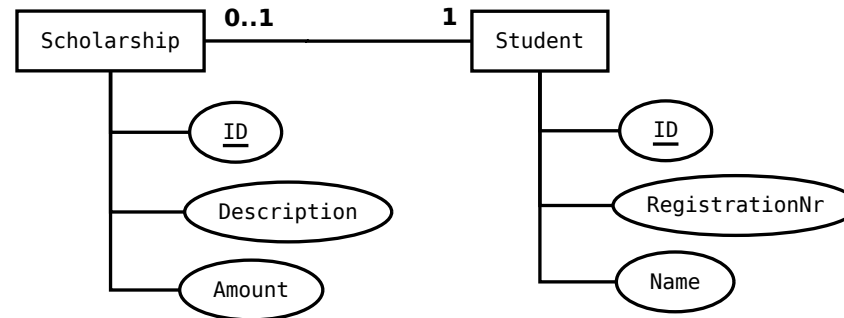
#### @JoinColumn

Specifies a column for joining an entity association or element collection.  
In this case: the name of the column, where the foreign key will be stored.



# Relationship mapping

Bidirectional OneToOne using foreign key - resulting SQL DDL

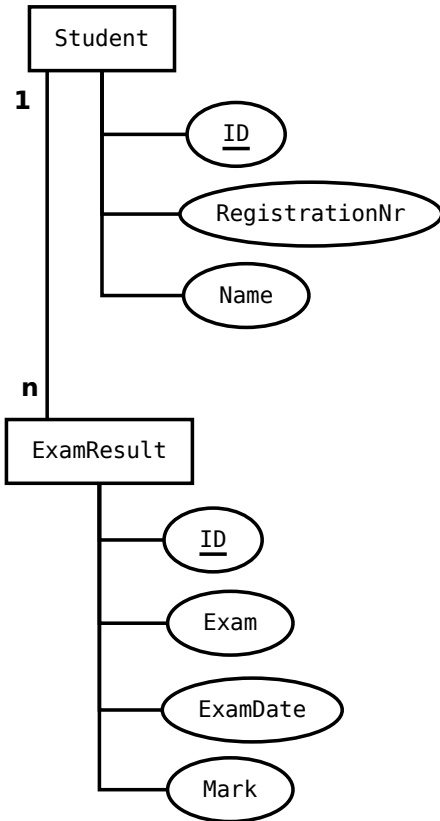


```
CREATE TABLE SCHOLARSHIP
(
    ID BIGINT PRIMARY KEY NOT NULL,
    AMOUNT INTEGER,
    DESCRIPTION VARCHAR(255),
    STUDENT_ID BIGINT,
    FOREIGN KEY ( STUDENT_ID )
    REFERENCES STUDENT ( ID )
);
CREATE UNIQUE INDEX uniqueIndexName
ON SCHOLARSHIP ( STUDENT_ID );
```

```
CREATE TABLE STUDENT
(
    ID BIGINT PRIMARY KEY NOT NULL,
    NAME VARCHAR(255),
    MATRIKELNUMMER VARCHAR(255)
);
CREATE UNIQUE INDEX anIndexName
ON STUDENT ( MATRIKELNUMMER );
```

# Relationship mapping

## Bidirectional OneToMany



The “non-owning” side

**@Entity**

```
public class Student extends BaseEntity {
```

```
    @Column(name = "matrikelNummer", unique = true)
```

```
    private String registrationNumber;
```

```
    private String name;
```

```
    @OneToOne(fetch = FetchType.LAZY, cascade = CascadeType.ALL,
              mappedBy = "grantedTo")
```

```
    private Scholarship scholarship;
```

```
    @OneToMany(cascade = CascadeType.ALL, mappedBy = "student")
```

```
    private List<ExamResult> examResults;
```

```
    @Transient
```

```
    private DateTime loginTime;
```

```
    ...
}
```

### Important annotations

**@OneToMany**


Defines a many-valued association with one-to-many multiplicity.



# Relationship mapping

## Bidirectional OneToMany cont'd

### The “owning” side



```
@Entity
public class ExamResult extends BaseEntity {

    @Column(name = "prufungsDatum")
    @Temporal(TemporalType.DATE)
    private Date examDate;

    private String exam;

    private int mark;

    @ManyToOne
    private Student student;

    @Transient
    private String examLocation;

    ...
}
```

### Important annotations


#### @ManyToOne

Defines a single-valued association to another entity class that has many-to-one multiplicity.



# Relationship mapping

## Bidirectional OneToMany - resulting SQL DDL




```
CREATE TABLE STUDENT
(
    ID BIGINT PRIMARY KEY NOT NULL,
    NAME VARCHAR(255),
    MATRIKELNUMMER VARCHAR(255)
);
CREATE UNIQUE INDEX anIndexNameB
ON STUDENT ( MATRIKELNUMMER );
```

```
CREATE TABLE EXAMRESULT
(
    ID BIGINT PRIMARY KEY NOT NULL,
    EXAM VARCHAR(255),
    PRUFUNGSDATUM DATE,
    MARK INTEGER NOT NULL,
    STUDENT_ID BIGINT,
    FOREIGN KEY ( STUDENT_ID ) REFERENCES STUDENT ( ID )
);
```

# Relationship mapping

## ManyToMany



### Important annotations

#### @ManyToMany

Defines a many-valued association with many-to-many multiplicity.

#### @Entity

```
public class Student extends BaseEntity {

    @Column(name = "matrikelNummer", unique = true)
    private String registrationNumber;

    private String name;

    @OneToOne(fetch = FetchType.LAZY,
              cascade = CascadeType.ALL,
              mappedBy = "grantedTo")
    private Scholarship scholarship;

    @OneToMany(cascade = CascadeType.ALL,
              mappedBy = "student")
    private List<ExamResult> examResults;


    @ManyToMany(mappedBy = "students")
    private List<Course> courses;

    @Transient
    private DateTime loginTime;

    ...
}
```

# Relationship mapping

## ManyToMany cont'd



```
@Entity
public class Course extends BaseEntity {

    private String courseNumber;


    private String title;

    @ManyToMany
    private List<Student> students;

    ...
}
```

# Relationship mapping

## ManyToMany - resulting SQL DDL



```
CREATE TABLE STUDENT
(
    ID BIGINT PRIMARY KEY NOT NULL,
    NAME VARCHAR(255),
    MATRIKELNUMMER VARCHAR(255)
);
CREATE UNIQUE INDEX anIndexNameB
ON STUDENT ( MATRIKELNUMMER );
```

```
CREATE TABLE COURSE
(
    ID BIGINT PRIMARY KEY NOT NULL,
    COURSENUMBER VARCHAR(255),
    TITLE VARCHAR(255)
);
```

```
CREATE TABLE COURSE_STUDENT
(
    COURSES_ID BIGINT NOT NULL,
    STUDENTS_ID BIGINT NOT NULL,
    FOREIGN KEY ( COURSES_ID )
    REFERENCES COURSE ( ID ),
    FOREIGN KEY ( STUDENTS_ID )
    REFERENCES STUDENT ( ID )
);
```



# Cascade and Fetch

---

- Cascade Types
  - All four relationship annotations may specify operations cascaded to associated entities
  - ALL, PERSIST, MERGE, REMOVE, REFRESH, DETACH
  - Default is none
- Orphan Removal
  - For @OneToOne and @OneToMany relationships
  - Default is false
- Fetching Strategies
  - Define how object hierarchies are loaded
  - EAGER = load all related objects immediately
  - LAZY = load the related objects only if they are accessed for the first time
  - Be careful with EAGER, as large object graphs may be loaded unintentionally!

# Persistence Concepts

---

## ■ Persistence Unit (PU)

- Defines a set of entity classes managed by the EntityManager instance in an application
- Maps the set of entity classes to a relational database

## ■ Persistence Context (PC)

- Set of managed entity instances that exist in a particular data store
- Runtime context

## ■ Entity Manager (EM)

- API for interaction with the persistence context
- Manipulates and controls the lifecycle of a persistence context
- Creates and removes persistent entity instances
- Finds entities using primary keys
- Runs queries on entities

# Persistence Unit

## ■ Persistence Unit

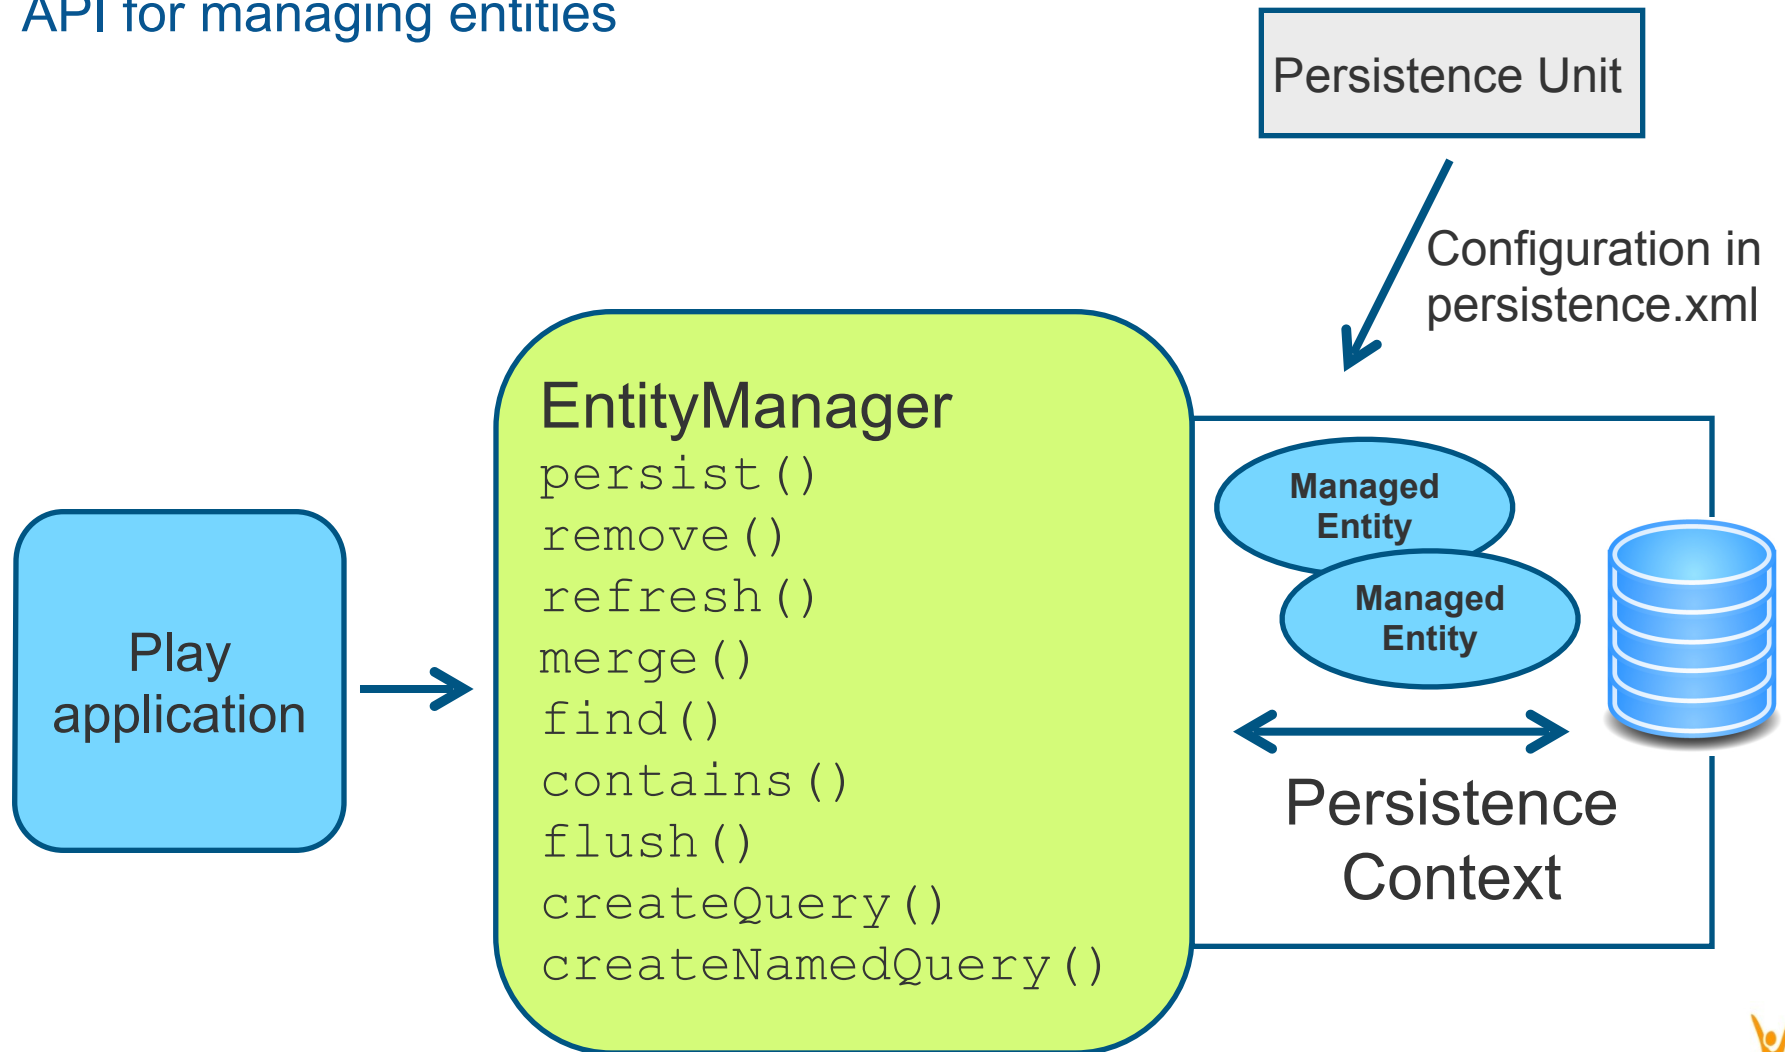
- Configuration to map entity classes in an application to a relational database
- `persistence.xml` defines one or more persistence units
  - Defined under `/conf/META-INF/persistence.xml`
  - Classes with JPA annotations are automatically detected upon start of the application

Names must be referenced in `application.conf`

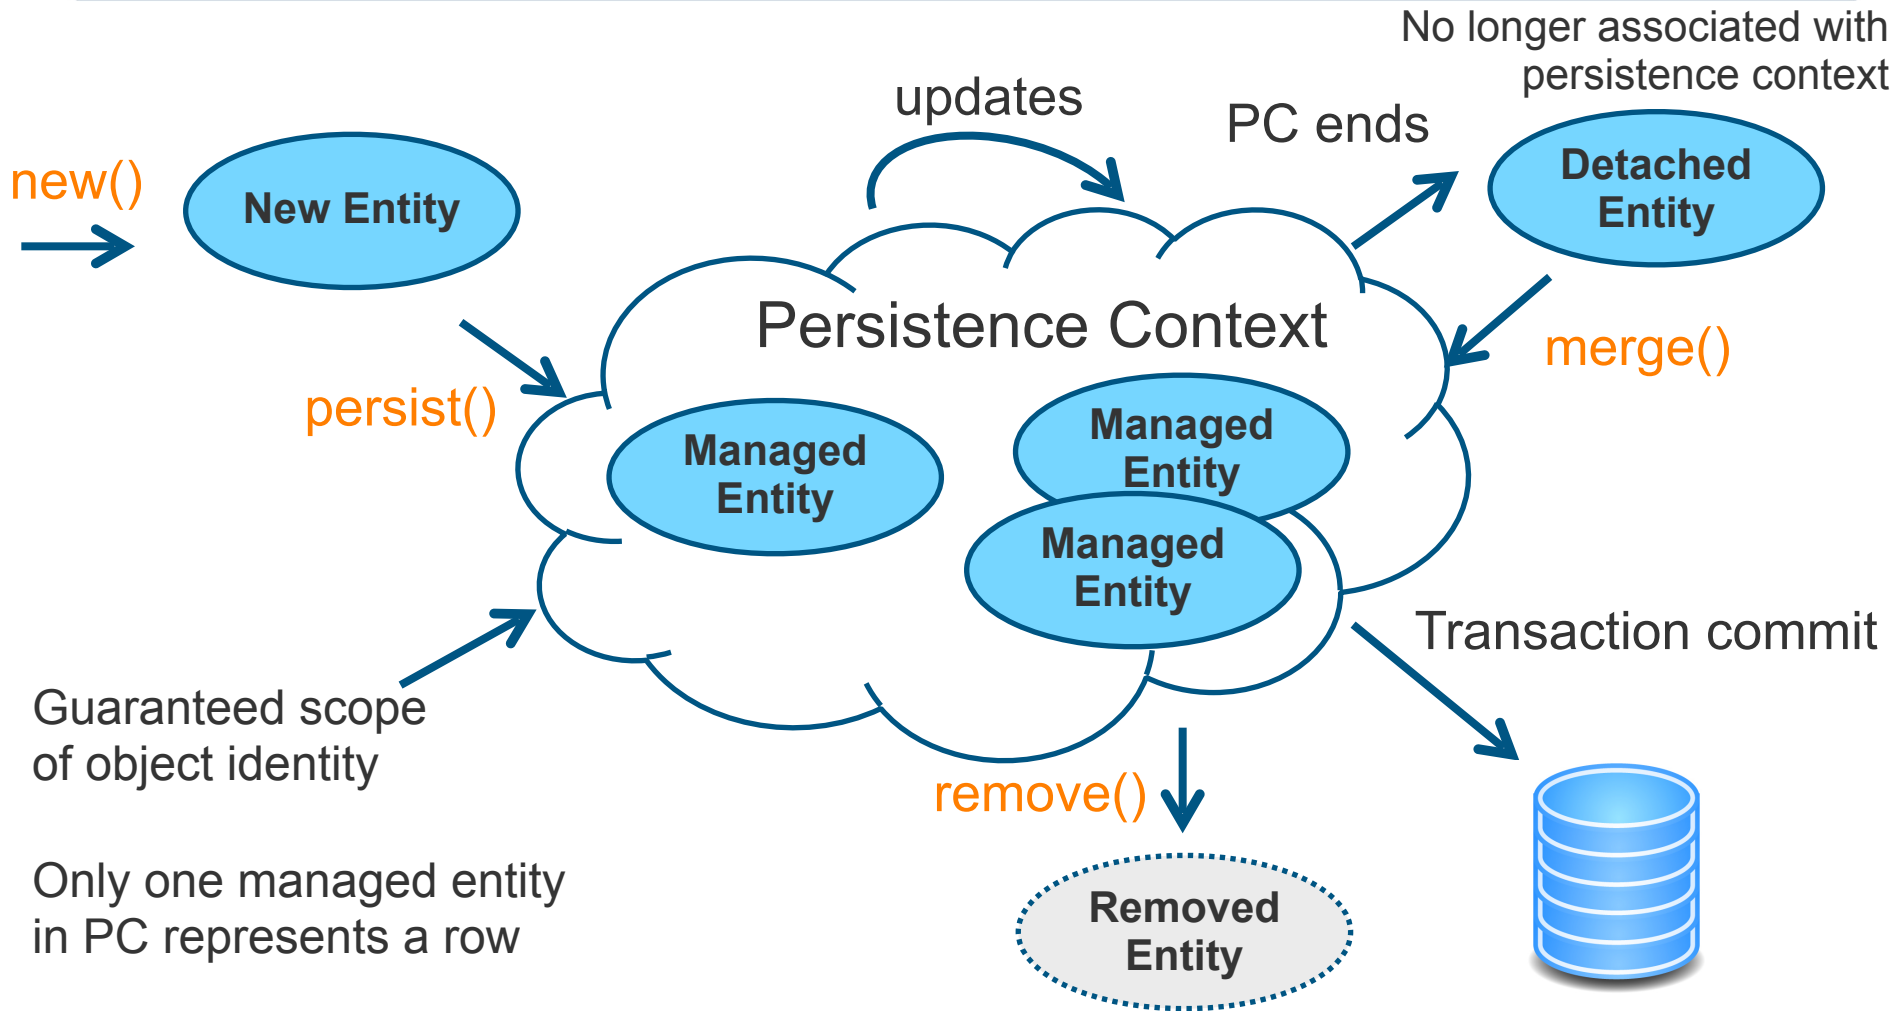
```
<persistence xmlns="http://java.sun.com/xml/ns/persistence"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xsi:schemaLocation="http://java.sun.com/xml/ns/persistence http://
java.sun.com/xml/ns/persistence/persistence_2_0.xsd"
  version="2.0">
  <persistence-unit name="defaultPersistenceUnit" transaction-type="RESOURCE_LOCAL">
    <provider>org.hibernate.ejb.HibernatePersistence</provider>
    <non-jta-data-source>DefaultDS</non-jta-data-source>
    <properties>
      <property name="hibernate.dialect" value="org.hibernate.dialect.H2Dialect"/>
      <property name="hibernate.hbm2ddl.auto" value="create" />
    </properties>
  </persistence-unit>
</persistence>
```

# Entity Manager

API for managing entities



# Entity lifecycle



Entities in managed/persistent state may be manipulated by the application and any changes will be **automatically detected and persisted** when the persistence context is flushed. There is no need to call a particular method to make your modifications persistent.

# Entity manager samples

Typically one addresses the entity manager from a dedicated persistence service class

---

```
public void persist(BaseEntity entity) {  
    em().persist(entity);  
}
```

Make an entity instance managed and persistent.  
Throws EntityExistsException, if the entity instance already exists.

```
public <T extends BaseEntity> T merge(T entity) {  
    return em().merge(entity);  
}
```

Make an entity instance managed and persistent. If it does not exist yet, persist it. If it already exists, the entity instance is updated.

```
public <T extends BaseEntity> T findEntity(Long id,  
                                           Class<T> entityClazz) {  
    return em().find(entityClazz, id);  
}
```

Find the entity instance using the primary key.

```
public void remove(BaseEntity entity) {  
    em().remove(entity);  
}
```

Remove the instance.

# Finding Entities

---

- Find entity by primary key using EntityManager

```
<T> T find(Class<T> entityClass, Object primaryKey)
```

- Example:

```
Student student = entityManager.find(Student.class, id);
```

- For complex queries use

- Java Persistence Query Language (JPQL) or Hibernate Query Language (HQL)
  - Both are object model focused query languages similar in nature to SQL. JPQL is a heavily-inspired-by subset of HQL. A JPQL query is always a valid HQL query, the reverse is not true however.
- Criteria API
- Native Queries

- EntityManager provides methods for creating Query objects

- createQuery
- createNativeQuery (using plain SQL - not recommended)

# JPQL/HQL

---

- Similar to SQL
- Works with entities as defined in the application and \*not\* with SQL table names and attribute names
- Portable (they abstract from vendor-specific SQL)
- Returns entities (no need to worry about result sets and their manual conversion to POJOs)
- Select, update, delete
  
- Support for
  - Joins
  - Conditional Expressions
  - Functional Expressions
  - Subqueries
  - Order by, group by, having
  - ...



# Querying Entities with JPQL

---

- Dynamic Query
  - Use parameter substitution and do not concatenate the JPQL string with the parameter values

```
public List<Student> getStudentByName(String studentName) {  
    TypedQuery<Student> studentTypedQuery =  
        em().createQuery("SELECT s FROM Student s WHERE s.name LIKE :studentName",  
            Student.class);  
  
    studentTypedQuery.setParameter("studentName", studentName);  
  
    return studentTypedQuery.getResultList();  
}
```

# Querying Entities with JPQL

---

- Static Query
  - Named Query
  - Recommended, as it may leverage use of query cache

```
@NamedQuery(name="findAllStudents",  
            query="SELECT s FROM Student s")  
  
@Entity  
public class Student extends BaseEntity {  
    ...  
}
```

```
public List<Student> getAllStudents() {  
    List<Student> students = em().createNamedQuery("findAllStudents",  
                                                Student.class).getResultList();  
    return students;  
}
```



# Criteria API

---

- Alternative to JPQL, same scope
- Dynamic Queries only
- Clauses are set using Java programming language objects
  - the query can be created in a typesafe manner
- Obtain a `CriteriaBuilder` instance by using the `EntityManager.getCriteriaBuilder` method

# Querying Entities with Criteria API

---

```
public Student getStudent(String registrationNumber) {  
  
    CriteriaBuilder cb = em().getCriteriaBuilder();  
    CriteriaQuery<Student> criteriaQuery = cb.createQuery(Student.class);  
    Root<Student> s = criteriaQuery.from(Student.class);  
    ParameterExpression<String> parameter = cb.parameter(String.class);  
    criteriaQuery.select(s).where(cb.equal(s.get("registrationNumber"), parameter));  
  
    TypedQuery<Student> typedQuery = em().createQuery(criteriaQuery);  
    typedQuery.setParameter(parameter, registrationNumber);  
    return typedQuery.getSingleResult();  
}
```

# Querying entities with Hibernate Criteria

---

```
public List<ExamResult> getNegativeExamResults(Student student) {  
    Criteria c = ((Session) JPA.em().getDelegate()).createCriteria(Student.class);  
    c.createCriteria("examResults").add(Restrictions.eq("mark", 5));  
    return c.list();  
}
```

# Wrap up

## Lessons learned today

---

- JPA/Hibernate provide a powerful ORM feature for Java-based applications
- Lot's of magic happens under the hood - know the data engineering basics first!
- Before putting your persistence layer into production, thoroughly test it using unit test



## References

---

1. Sun Microsystems. JSR 220: Enterprise JavaBeans™, Version 3.0 – Java Persistence API, 2006
2. Carol McDonald. Java Persistence API: Best Practices, Sun Tech Days 2008-2009  
<http://de.slideshare.net/caroljmcDonald/td09jpabestpractices2>
3. Carol McDonald. Enterprise JavaBean 3.0 & Java Persistence APIs: Simplifying Persistence, Sun Tech Days 2006-2007  
<http://de.slideshare.net/caroljmcDonald/persistenceecmcdonaldmainejug3>
4. The Java EE 6 Tutorial, <http://docs.oracle.com/javaee/6/tutorial/doc/bnbpy.html>
5. <https://www.codecentric.de/files/2011/05/flush-und-clear-or-mapping-anti-patterns.pdf> (in German)

# Backup

Further background information for the interested

---



# Object relational mapping (ORM)

Why objects and databases do not play well together

---

- **Object-Relational Impedance Mismatch (or paradigm mismatch)**
  - RDBMS represent data in tabular format
  - Object-oriented languages such as Java present data in an interconnected graph of objects
- **Loading and storing objects using a tabular relational database exposes different problems:**
  - **Granularity**

Oftentimes the object model will contain more classes, than the number of corresponding tables in the database
  - **Subtypes**

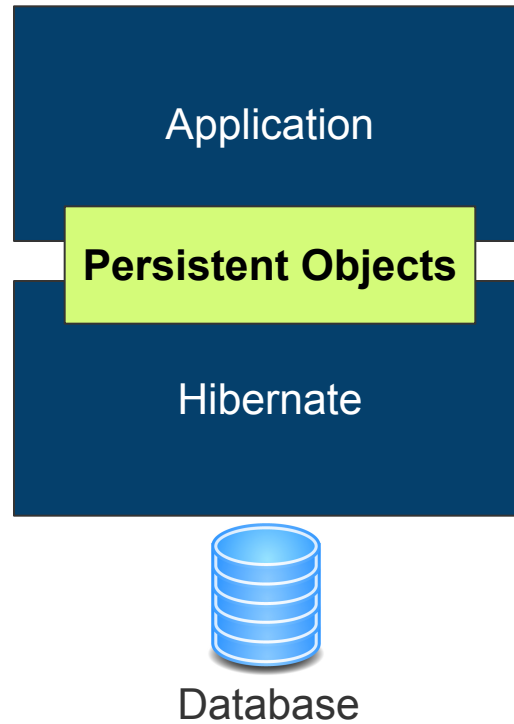
Inheritance is an integral part of object-oriented programming. RDBMS usually do not foresee an inheritance mechanism.
  - **Identity**

A RDMS defines a single notion of sameness: the primary key. Java, however, defines both, object identity `a==b` and object equality `a.equals(b)`
  - **Associations**

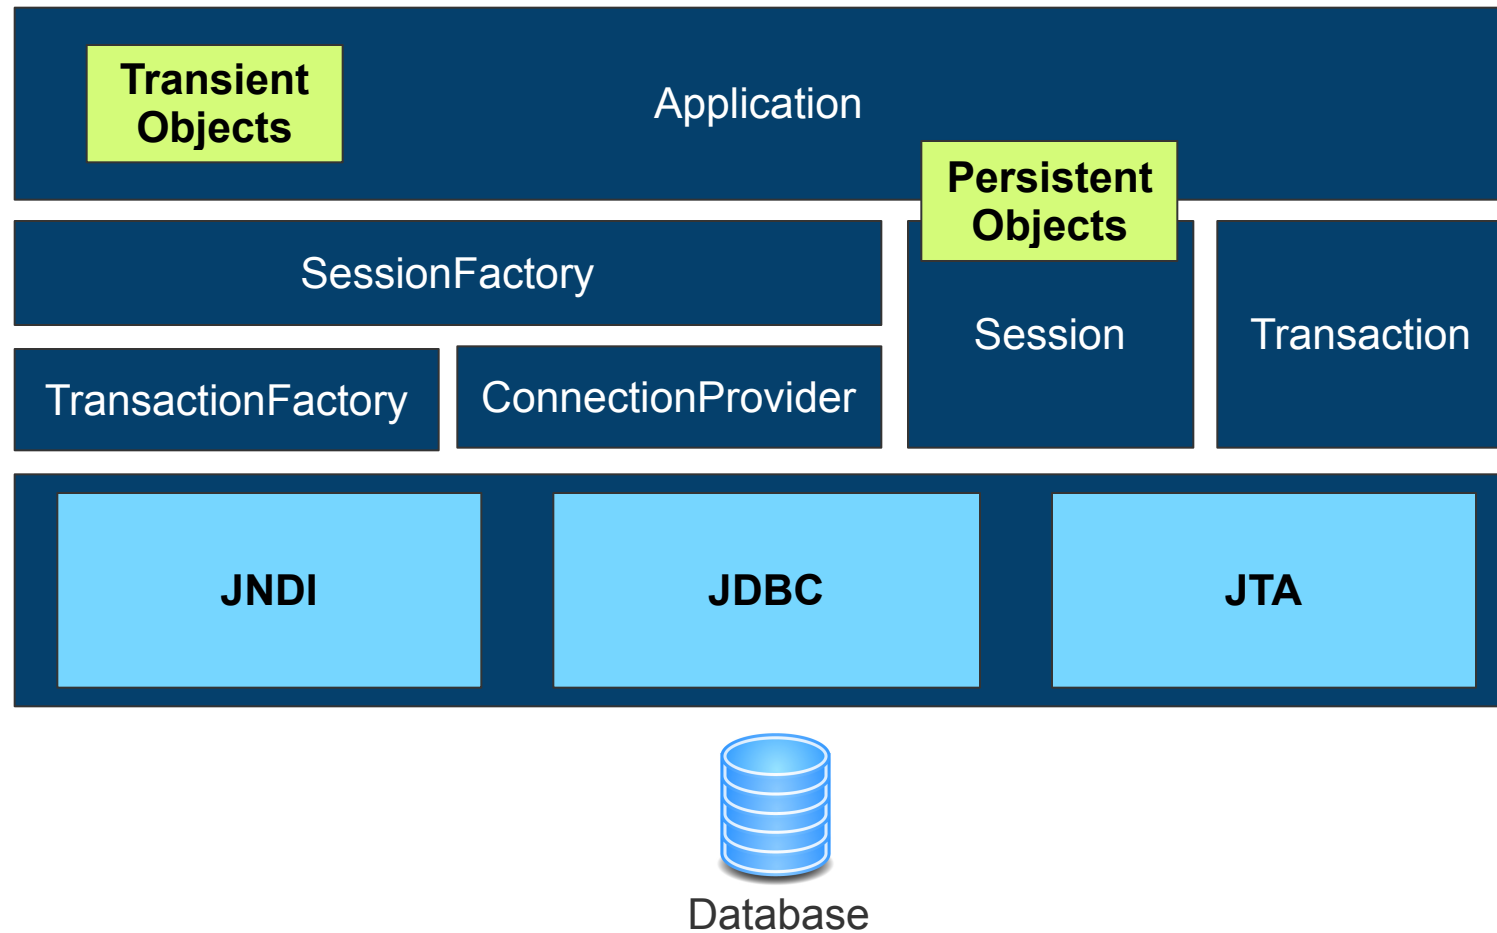
Associations are represented as unidirectional references in an Object Oriented Language such as Java. An RDMS uses the concept of foreign keys. If one requires bidirectional relationships in Java, an association must be defined twice.
  - **Data navigation.** Association style navigation (Java), vs. SQL JOINS.

# High-level overview of the Hibernate architecture

---



# Detailed view of the Hibernate architecture



# Transaction

What exactly makes a database transaction?

---

- A transaction is a sequence of operations, performed as a single logical unit of work. The single logical unit of work must have four properties in order to qualify it as a transaction.
  - **Atomicity**  
A transaction must be an atomic unit of work, i.e., all of its data modifications are performed, or no modification is performed at all.
  - **Consistency**  
After a transaction is completed, all data must be left in a consistent state. The written data must confirm to the defined rules such as constraints, triggers, cascades, etc.  
All internal data structures, such as indexes, must be correct at the end of the transaction.



# Transaction cont'd

What exactly makes a database transaction?

---

- **Isolation**

Modifications of a given transaction must be isolated from modifications made by other concurrent transactions. A transaction never recognizes data in an intermediate state, which was potentially caused by another concurrent transaction.

- **Durability**

After a transaction has been completed, its effects are permanently stored in the system. Modifications persist even in the case of a system failure.